

Palette: Image-to-Image Diffusion Models

Chitwan Saharia, William Chan, Huiwen Chang, Chris A. Lee, Jonathan Ho, Tim Salimans,
David J. Fleet* and **Mohammad Norouzi***
{sahariac, williamchan, davidfleet, mnorouzi}@google.com
Google Research

Abstract

We introduce Palette, a simple and general framework for image-to-image translation using conditional diffusion models. On four challenging image-to-image translation tasks (colorization, inpainting, uncropping, and JPEG decompression), Palette outperforms strong GAN and regression baselines, and establishes a new state of the art. This is accomplished without task-specific hyper-parameter tuning, architecture customization, or any auxiliary loss, demonstrating a desirable degree of generality and flexibility. We uncover the impact of using L_2 vs. L_1 loss in the denoising diffusion objective on sample diversity, and demonstrate the importance of self-attention through empirical architecture studies. Importantly, we advocate a unified evaluation protocol based on ImageNet, and report several sample quality scores including FID, Inception Score, Classification Accuracy of a pre-trained ResNet-50, and Perceptual Distance against reference images for various baselines. We expect this standardized evaluation protocol to play a critical role in advancing image-to-image translation research. Finally, we show that a single generalist Palette model trained on 3 tasks (colorization, inpainting, JPEG decompression) performs as well or better than task-specific specialist counterparts. Check out <https://bit.ly/palette-diffusion> for more details.

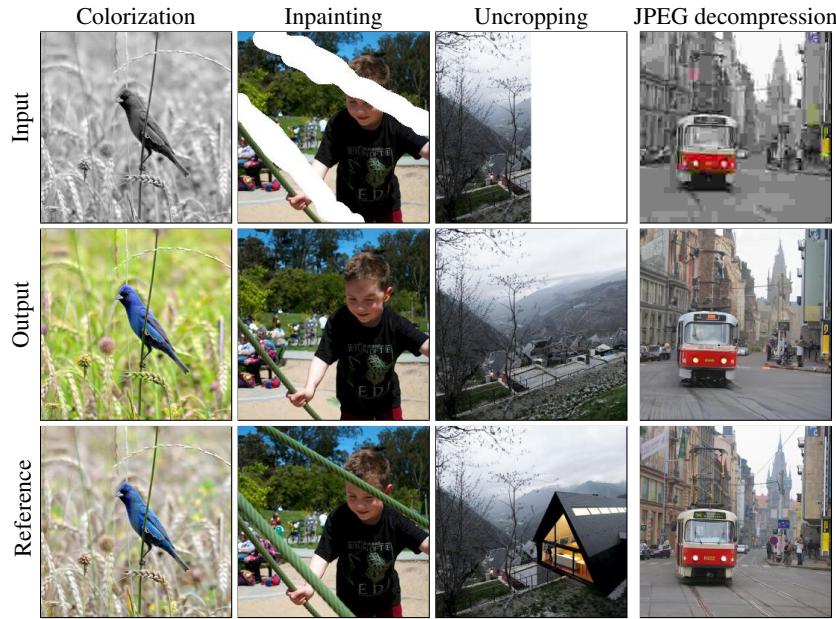


Figure 1: Illustration of Palette’s performance on four image-to-image translation tasks.

*Equal Contribution



Figure 2: Palette panorama uncropping. Given the center 256×256 pixels, we extrapolate 2 times to the left (4 applications of 50% Palette uncropping), and 2 times to the right, to generate the final 256×1280 panorama. See Figure C.10 in the Appendix for more samples.

1 Introduction

Many problems in computer vision and image processing can be formulated as image-to-image translation. Examples include restoration tasks, such as super-resolution, colorization, and inpainting, as well as pixel-level image understanding tasks, such as instance segmentation and the estimation of intrinsic images. Many such tasks, like those in Figure 1, are complex inverse problems, where multiple output images are consistent with a single input. A natural approach to image-to-image translation is to learn the conditional distribution of output images given the input using deep generative models that can capture multi-modal distributions in the high-dimensional space of images. Generative Adversarial Networks (GANs) [Goodfellow et al., 2014, Radford et al., 2015] have emerged as the model family of choice for many image-to-image translation tasks [*e.g.*, Isola et al., 2017a], as they are capable of generating high fidelity outputs, are broadly applicable, and support efficient sampling. Nevertheless, GANs can be challenging to train [*e.g.*, Arjovsky et al., 2017, Gulrajani et al., 2017], and often drop modes in the output distribution [*e.g.*, Metz et al., 2016, Ravuri and Vinyals, 2019]. Other generative models such as Autoregressive Models [van den Oord et al., 2016, Parmar et al., 2018], VAEs [Kingma and Welling, 2013, Vahdat and Kautz, 2020], and Normalizing Flows [Dinh et al., 2016, Kingma and Dhariwal, 2018] have seen success in specific applications, but arguably, have not established the same level of sample quality and general applicability as GANs.

Diffusion models [Sohl-Dickstein et al., 2015, Song and Ermon, 2020, Ho et al., 2020] have received a surge of recent interest [*e.g.*, Cai et al., 2020, Song et al., 2021, Hoogeboom et al., 2021, Vahdat et al., 2021, Kingma et al., 2021, Austin et al., 2021], resulting in several key advances in modeling continuous data. On speech synthesis, diffusion models have achieved human evaluation scores on par with SoTA Autoregressive Models [Chen et al., 2021a,b, Kong et al., 2021]. On the class-conditional ImageNet generation challenge, they have outperformed the strongest GAN baselines in terms of FID scores [Dhariwal and Nichol, 2021, Ho et al., 2021]. On image super-resolution, they have delivered impressive face enhancement results, outperforming GANs [Saharia et al., 2021a]. Despite these results, it is not clear whether diffusion models can present a versatile and general solution to the problem of image-to-image translation, rivaling GANs.

This paper investigates the general applicability of *Palette*, our implementation of image-to-image diffusion models, to a suite of four distinct and challenging tasks, namely colorization, inpainting, uncropping (*a.k.a* outpainting or extrapolation), and JPEG artifact removal (see Figure 1 and Figure 2). Our empirical results suggest that *Palette* with no task-specific architecture customization, and no changes to the hyper-parameters or the loss, can deliver high-fidelity outputs across all four tasks. It outperforms several existing baselines and our own strong regression baseline, and thereby establishes a new state-of-the-art. Importantly, we show that a single *generalist* *Palette* model, trained on colorization, inpainting and JPEG decompression, outperforms a task-specific JPEG decompression variant and achieves competitive performance on the other tasks.

We study key components of *Palette*, including the loss function and the neural net architecture. We find that while L_2 [Ho et al., 2020] and L_1 [Chen et al., 2021a] losses in the denoising objective yield similar sample-quality scores, L_2 leads to a higher degree of diversity in model samples, whereas L_1 [Chen et al., 2021a] produces more conservative outputs. We also find that removing self-attention layers from the U-Net architecture of *Palette*, in order to build a fully convolutional model, hurts performance. We present a set of evaluation protocols for inpainting, uncropping, and JPEG decompression based on ImageNet [Deng et al., 2009]. We report several quantitative evaluation metrics, and contribute a surprisingly strong regression baseline. We expect this evaluation protocol to play a critical role in advancing image-to-image translation research. Our key contributions include:

1. We study and confirm the versatility and general applicability of diffusion models to image-to-image translation. Palette, our implementation of image-to-image diffusion models, is able to achieve SoTA performance on four tasks, colorization, inpainting, uncropping, and JPEG decompression, with no task-specific tuning or customization.
2. We recognize the impact of using L_1 vs. L_2 in the denoising objective on sample diversity, and show the importance of self-attention through empirical studies of the neural network architecture.
3. We propose a unified evaluation protocol across image translation tasks based on the ImageNet dataset, and report several sample quality scores for several baselines including Palette, a surprisingly strong regression baseline, and prior work where possible.
4. We show that a multi-task Palette model performs as well or better than task-specific models.

2 Related work

Our work is inspired by Pix2Pix [Isola et al., 2017a], which explored a diverse array of image-to-image translation tasks for GAN based models. Other GAN-based techniques have been proposed for various image-to-image translation problems such as unpaired translation [Zhu et al., 2017], unsupervised cross-domain generation [Taigman et al., 2016], multi-domain translation [Choi et al., 2018], few shot translation [Liu et al., 2019] and many more. Nevertheless, existing GAN models are sometimes unsuccessful in holistically translating images with consistent structural and textural regularity. Diffusion models [Sohl-Dickstein et al., 2015] recently emerged with impressive results on image generation [Ho et al., 2020, 2021, Dhariwal and Nichol, 2021], audio synthesis [Chen et al., 2021a, Kong et al., 2020], and image super-resolution [Saharia et al., 2021a, Kadkhodaie and Simoncelli, 2021], as well as unpaired image-to-image translation [Sasaki et al., 2021] and image editing [Meng et al., 2021, Sinha et al., 2021]. Our conditional diffusion models build on these recent advances, showing versatility and generality on a suite of image-to-image translation tasks.

Early **inpainting** approaches [*e.g.*, Bertalmio et al., 2000, Barnes et al., 2009, He and Sun, 2012, Hays and Efros, 2007, Roth and Black, 2005] work well on textured regions but often fall short in generating semantically consistent structure. GAN models are now widely used, but often require auxiliary objectives on structures, context, edges, contours and hand-engineered features [*e.g.*, Iizuka et al., 2017, Yu et al., 2018a, 2019, Nazeri et al., 2019, Yi et al., 2020, Liu et al., 2020, Kim et al., 2021], and a lack of stochasticity and diversity in their outputs has been observed [Zheng et al., 2019, Zhao et al., 2021]. Generic diffusion models for image generation can be used for linear inverse tasks like inpainting [Sohl-Dickstein et al., 2015, Song et al., 2020, Kadkhodaie and Simoncelli, 2021], but we posit that conditional models trained for inpainting will be superior. **Image uncropping** (*a.k.a.* outpainting) is considered more challenging than inpainting as it entails generating open-ended content with less context. Early methods relied on retrieval [Kopf et al., 2012, Wang et al., 2014, Shan et al., 2014]. GAN-based methods are now predominant [*e.g.*, Teterwak et al., 2019], but are often domain-specific [*e.g.*, Yang et al., 2019, Bowen et al., 2021, Wang et al., 2019a, Cheng et al., 2021, Lin et al., 2021]. We show that conditional diffusion models trained on large datasets reliably address both inpainting and uncropping on a wide range of image domains.

Colorization is a well-studied image-to-image task [*e.g.*, Kumar et al., 2021, Guadarrama et al., 2017, Royer et al., 2017, Ardizzone et al., 2019], requiring a degree of scene understanding, which makes it a natural choice for self-supervised learning [Larsson et al., 2016]. Challenges include diverse colorization [*e.g.*, Deshpande et al., 2017], respecting semantic categories [*e.g.*, Zhang et al., 2016], and producing high-fidelity color [*e.g.*, Guadarrama et al., 2017]. While some prior work makes use of specialized auxiliary classification losses, we find that generic image-to-image diffusion models work well without task-specific specialization. **JPEG decompression** is a nonlinear inverse problem for restoring realistic textures and removing compression artifacts. Dong et al. [2015] applied deep CNN architectures for JPEG decompression. Galteri et al. [2017, 2019] have successfully applied GAN models for decompression tasks, but they have been restricted to quality factors of above 10. We show the effectiveness of Palette in removing compression artifacts for quality factors as low as 5.

Multi-task training is a relatively under-explored area in image-to-image translation. A few papers [*e.g.*, Qian et al., 2019, Yu et al., 2018b] have focused on simultaneous training over multiple tasks, but they are primarily focused on enhancement tasks like deblurring, denoising, and super-resolution, and they use smaller, modular networks. Several works have also dealt with simultaneous training over multiple degradations on a single task *e.g.*, multi-scale super-resolution [Kim et al., 2016], jpeg-decompression on multiple quality factors [Galteri et al., 2019, Liu et al., 2018]. This paper takes the first step toward building multi-task image-to-image diffusion models for a variety of tasks.

3 Palette

Background on Conditional Diffusion Models. Diffusion models [Sohl-Dickstein et al., 2015, Ho et al., 2020] convert samples from a standard Gaussian distribution into samples from an empirical data distribution through an iterative denoising process. Conditional diffusion models [*e.g.*, Chen et al., 2021a, Saharia et al., 2021b] make the denoising process conditional on an input signal. Image-to-image diffusion models are conditional diffusion models of the form $p(\mathbf{y} \mid \mathbf{x})$, where both \mathbf{x} and \mathbf{y} are images, *e.g.*, \mathbf{x} is a grayscale image and \mathbf{y} is a reference color image. These models have been previously applied to image super-resolution [Saharia et al., 2021a, Nichol and Dhariwal, 2021]. We study the general applicability of image-to-image diffusion models on a broad set of tasks.

Loss function. For a detailed treatment of diffusion models, please see Appendix A. Here, we briefly discuss the denoising loss function. Given a reference output image \mathbf{y} , we generate a noisy version $\tilde{\mathbf{y}}$, and train a neural network f_θ to denoise $\tilde{\mathbf{y}}$ given \mathbf{x} and a noise level indicator γ . The objective takes the form

$$\mathbb{E}_{(\mathbf{x}, \mathbf{y})} \mathbb{E}_{\epsilon \sim \mathcal{N}(0, I)} \mathbb{E}_\gamma \left\| f_\theta(\mathbf{x}, \underbrace{\sqrt{\gamma} \mathbf{y} + \sqrt{1 - \gamma} \epsilon, \gamma}_{\tilde{\mathbf{y}}} - \epsilon \right\|_p^p, \quad (1)$$

Chen et al. [2021a] and Saharia et al. [2021a] suggest using the L_1 norm, *i.e.*, $p = 1$, whereas the standard formulation is based on the usual L_2 norm [*e.g.*, Ho et al., 2020]. We perform careful ablations in what follows, and recognize the impact of the choice of L_p on sample diversity. That is, L_1 yields significantly lower sample diversity. While L_1 may be useful, to reduce potential hallucinations in some applications, here we adopt L_2 to capture the output distribution more faithfully.

Architecture. Palette uses a standard U-Net architecture [Ho et al., 2020], with several architecture modifications inspired by recent work [Song et al., 2021, Saharia et al., 2021a, Dhariwal and Nichol, 2021]. The network architecture is based on the 256×256 class-conditional U-Net model of Dhariwal and Nichol [2021]. The two main differences between our architecture and theirs are (i) absence of class-conditioning, and (ii) additional conditioning of the source image through concatenation, similar to Saharia et al. [2021a].

4 Evaluating Image-to-image Translation Models

Evaluating image-to-image translation models is challenging. Prior work on colorization [Zhang et al., 2016, Guadarrama et al., 2017, Kumar et al., 2021] has relied on FID scores and human evaluation for model comparison. Tasks like inpainting [Yu et al., 2019, 2018a] and uncropping [Teterwak et al., 2019, Wang et al., 2019b] have often heavily relied on qualitative evaluation. For many tasks, such as JPEG artifact removal [Dong et al., 2015, Liu et al., 2018, Galteri et al., 2019], it has been common to use reference-based pixel-level similarity scores such as PSNR and SSIM. It is also notable that many tasks lack a standardized dataset for evaluation, *e.g.*, different test sets with method-specific splits are used for evaluation.

We propose a unified evaluation protocol for inpainting, uncropping, and JPEG decompression on ImageNet [Deng et al., 2009], due to its scale, diversity, and public availability. For inpainting and uncropping, existing work has primarily relied on Places2 dataset [Zhou et al., 2017] for evaluation. Hence, we also use a standard evaluation setup on Places2 for these tasks. Specifically, we advocate the use of ImageNet ctest10k split proposed by Larsson et al. [2016] as a standard subset for benchmarking of all image-to-image translation tasks on ImageNet. We also introduce a similar category balanced 10,950 image subset of Places2 validation set called *places10k*. We further advocate the use of automated metrics that can capture both image quality and diversity, in addition to controlled human evaluation. We avoid the use of pixel-level metrics such as PSNR and SSIM since such metrics are not reliable measures of sample quality for difficult tasks that require hallucination; *i.e.*, recent super-resolution works [*e.g.*, Ledig et al., 2017, Dahl et al., 2017, Menon et al., 2020] observe that PSNR and SSIM do not correlate well with human judgement as they prefer blurry regression outputs.

We use four quantitative measures of sample quality for image-to-image translation: **Inception Score (IS)** [Salimans et al., 2017]; **Fréchet Inception Distance (FID)**; **Classification Accuracy (CA)** (top-1) of a pre-trained ResNet-50 classifier; and a simple measure of **Perceptual Distance (PD)**, *i.e.*, Euclidean distance in the Inception-v1 feature space. We also assess **sample diversity** by visual inspection and histogram plots of pairwise SSIM scores between multiple model outputs. Sample diversity is challenging and has been a key limitation of many existing GAN-based methods.

The ultimate evaluation of image-to-image translation models is **human evaluation**; *i.e.*, whether or not humans can discriminate model outputs from reference images. To this end we use 2-alternative forced choice (2AFC) trials to evaluate the perceptual quality of model outputs against reference images (*c.f.*, the Colorization Turing Test [Zhang et al., 2016]). We summarize the results in terms of the **fool rate**, the percentage of human raters who select model outputs over reference images when they were asked “Which image would you guess is from a camera?”. (See Appendix C for details.)

5 Experiments

We study the general applicability of *Palette* to a suite of four distinct and challenging image-to-image translation tasks:

1. **Colorization** transforms an input grayscale image to a plausible color image.
2. **Inpainting** fills in user-specified masked regions of an image with realistic content.
3. **Uncropping** extends an input image along one or more directions to enlarge the image.
4. **JPEG decompression** corrects for JPEG compression artifacts to recover plausible details.

We evaluate *Palette* on these tasks without task-specific hyper-parameter tuning, or architecture customization, or any auxiliary loss function. Inputs and outputs for all tasks are presented as RGB 256×256 images. Each of these tasks presents its own unique challenges (often requiring task-specific engineering), while enabling new applications. Colorization entails some representation of objects, segmentation and layout, often with long-range image dependencies. Inpainting is challenging with large masks and diverse image datasets comprising cluttered scenes. Outpainting is widely considered even more challenging than inpainting as there is less surrounding context to constrain semantically meaningful generation. While the other tasks are linear inverse problems, JPEG decompression is a non-linear inverse problem; it entails the separation of local signal properties from compression artifacts, thereby requiring a good model of local image statistics. While previous work has studied these problems extensively, it is rarely the case that a single model achieves SoTA performance on all tasks. In what follows, *Palette* uses an L_2 loss for the denoising objective, unless otherwise specified. More implementation details can be found in Appendix B.

5.1 Colorization

Prior work [*e.g.*, Zhang et al., 2016, Kumar et al., 2021] has adopted LAB or YCbCr color spaces to represent output images for colorization, but we use the RGB color space to maintain generality across tasks. Our preliminary experiments suggest that *Palette* is equally effective in YCbCr and RGB spaces. We compare *Palette* with cGAN [Isola et al., 2017b], PixColor [Guadarrama et al., 2017], and Coltran [Kumar et al., 2021]. We show example qualitative results in Figure 3 and quantitative scores in Table 1. *Palette* establishes a new SoTA, outperforming existing work by a large margin in sample quality and human evaluation scores. Surprisingly, our simple L_2 Regression baseline is also able to outperform prior techniques, highlighting the importance of modern architectures and large-scale training, even for a basic Regression model. *Palette* improves upon human raters’ fool rate of Coltran by more than 10%, approaching an ideal fool rate of 50%.



Figure 3: Illustration of colorization methods on ImageNet validation images. Baselines: [†][Guadarrama et al., 2017], [‡][Kumar et al., 2021], and our own strong regression baseline. Figure C.2 shows more samples.

Model	FID-5K ↓	IS ↑	CA ↑	PD ↓	Fool rate (3 sec) ↑
<i>Prior Work</i>					
cGAN [Isola et al., 2017b]	24.41	-	-	-	-
PixColor [Guadarrama et al., 2017]	24.32	-	-	-	29.90%
Coltran [Kumar et al., 2021]	19.37	-	-	-	36.55%
<i>This paper</i>					
Regression	17.89	169.8	68.2%	60.0	39.45%
Palette	15.78	200.8	72.5%	46.2	47.80%
Reference	14.68	229.6	75.6%	0.0	-

Table 1: Colorization quantitative scores and fool rates on ImageNet val set. Appendix C.1 has more results.

5.2 Inpainting

We follow Yu et al. [2019] and train inpainting models on free-form generated masks, augmented with simple rectangular masks. To maintain generality of Palette across tasks, in contrast to prior work, we do not pass a binary inpainting mask to the models. Instead, we fill the masked region with standard Gaussian noise, which is compatible with denoising diffusion models. We train the model to predict only the masked out pixels, instead of the entire image to speed up training. We compare Palette with DeepFillv2 [Yu et al., 2019], HiFill [Yi et al., 2020], and Photoshop’s *Content-aware Fill*, showing qualitative results in Figure 4 and quantitative scores in Table 2. Palette exhibits strong performance across inpainting datasets and mask configurations, establishing a new SoTA.

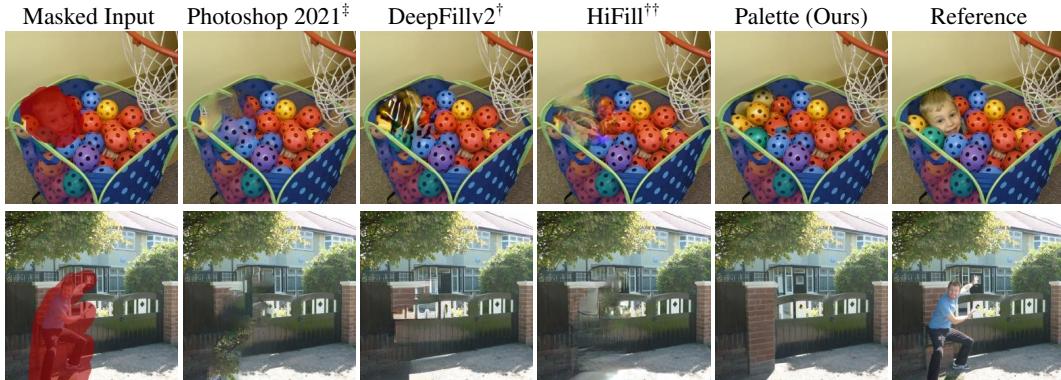


Figure 4: Comparison of inpainting methods on object removal. Baselines: [‡]Photoshop’s *Content-aware Fill* built on PatchMatch [Barnes et al., 2009], [†][Yu et al., 2019], and ^{††}[Yi et al., 2020]. Figure C.4 has more samples.

Mask Type	Model	ImageNet				Places2	
		FID ↓	IS ↑	CA ↑	PD ↓	FID ↓	PD ↓
20-30% <i>free form</i>	DeepFillv2 [Yu et al., 2019]	9.4	174.6	68.8%	64.7	13.5	63.0
	HiFill [Yi et al., 2020]	12.4	157.0	65.7%	86.2	15.7	92.8
	Palette (Ours)	5.2	205.5	72.3%	27.6	11.7	35.0
128×128 <i>center mask</i>	DeepFillv2 [Yu et al., 2019]	18.0	135.3	64.3%	117.2	15.3	96.3
	HiFill [Yi et al., 2020]	20.1	126.8	62.3%	129.7	16.9	115.4
	Palette (Ours)	6.6	173.9	69.3%	59.5	11.9	57.3
Reference		5.1	231.6	74.6%	0.0	11.4	0.0

Table 2: Quantitative evaluation for free-form and center inpainting on ImageNet and Places2 validation images.

5.3 Uncropping

Recent work [Teterwak et al., 2019, Lin et al., 2021] has shown impressive visual effects by extending (extrapolating) input images along the right border. We train Palette on uncropping in any one of

the four directions, or around the entire image border on all four sides. In all cases, we keep the area of the masked out region at 50% of the image. Like inpainting, we fill the masked out region with Gaussian noise, and keep the unmasked region fixed during inference. We compare Palette with Boundless [Teterwak et al., 2019] and InfinityGAN [Lin et al., 2021], showing qualitative results in Figure 5 and quantitative sample quality scores in Table 3. Palette outperforms existing baselines on both ImageNet and Places2 by a large margin, demonstrating the strong capability of Palette models in image uncropping without any task-specific modifications to the architecture or loss. Palette achieves significantly higher fool rates over Boundless, indicating more realistic looking outputs.



Figure 5: Image uncropping results on Places2 validation images. Baselines: Boundless[†] [Teterwak et al., 2019] and InfinityGAN^{‡‡} [Lin et al., 2021] trained on a scenery subset of Places2. Figure C.7 shows more samples.

Model	ImageNet				Places2		
	FID ↓	IS ↑	CA ↑	PD ↓	FID ↓	PD ↓	Fool rate ↑
Boundless [Teterwak et al., 2019]	18.7	104.1	58.8%	127.9	11.8	129.3	20.7%
Palette (Ours)	5.8	138.1	63.4%	85.9	3.53	103.3	39.9%
Reference	2.7	250.1	76.0%	0.0	2.1	0.0	-

Table 3: Quantitative scores and human raters’ Fool rates on uncropping.

5.4 JPEG Decompression

Finally, we evaluate Palette on the task of removing JPEG compression artifacts, a long standing image enhancement problem [Dong et al., 2015, Galteri et al., 2019, Liu et al., 2018]. Similar to prior work [Ehrlich et al., 2020, Liu et al., 2018], we train Palette on inputs compressed with various quality factors (QF). Prior work has typically limited itself to a Quality Factor higher than 10, we increase the difficulty of the problem and train on Quality Factors as low as 5, with more severe compression artifacts. Table 4 summarizes the results of Palette on ImageNet. Palette exhibits strong performance across all quality factors, being significantly better than the Regression baseline. As expected, the performance gap between Palette and Regression baseline widens with decreasing quality factor. Figure 6 shows the qualitative comparison between Palette and our Regression baseline at a quality factor of 5. It is easy to see that the regression model produces blurry outputs, while Palette produces sharper images with realistic textures.

5.5 The Role of Self-Attention in Diffusion Model Architectures

Self-attention layers [Vaswani et al., 2017] have been an important component in recent U-Net architectures for diffusion models [Ho et al., 2020, Dhariwal and Nichol, 2021]. While self-attention layers build a direct form of global dependency into the model, they prevent generalization to unseen image resolutions. Generalization to new resolutions at test time is convenient for many image-to-image translation tasks, for which previous works on image-to-image translation have primarily relied on fully convolutional architectures [Yu et al., 2019, Galteri et al., 2019].

We analyze the impact of these self-attention layers on sample quality for inpainting, one of the more difficult image-to-image translation tasks. In order to enable input resolution generalization for

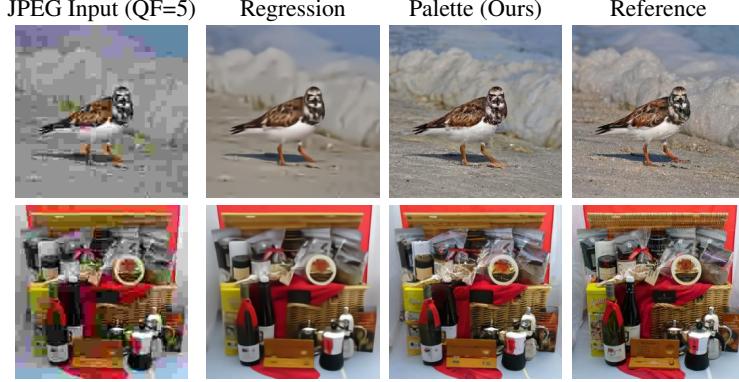


Figure 6: Visual comparison for JPEG decompression on ImageNet validation images.

QF	Model	FID-5K ↓	IS ↑	CA ↑	PD ↓
5	Regression	29.0	73.9	52.8%	155.4
	Palette (Ours)	8.3	133.6	64.2%	95.5
10	Regression	18.0	117.2	63.5%	102.2
	Palette (Ours)	5.4	180.5	70.7%	58.3
20	Regression	11.5	158.7	69.7%	65.4
	Palette (Ours)	4.3	208.7	73.5%	37.1
Ground Truth		2.7	250.1	75.9%	0.0

Table 4: Quantitative evaluation for JPEG decompression for various Quality Factors (QF).

Palette, we explore different alternatives to global self-attention layers in the U-Net architecture. In particular, we experiment with the following four configurations:

1. **Global Self-Attention:** Baseline configuration with global self-attention layers at 32×32 , 16×16 and 8×8 resolutions.
2. **Local Self-Attention:** Local self-attention layers [Vaswani et al., 2021] at 32×32 , 16×16 and 8×8 resolutions, at which feature maps are divided into 4 non-overlapping query blocks.
3. **More ResNet Blocks without Self-Attention:** $2 \times$ residual blocks at 32×32 , 16×16 and 8×8 resolutions allowing deeper convolutions to increase receptive field sizes.
4. **Dilated Convolutions without Self-Attention:** Similar to 3. ResNet blocks at 32×32 , 16×16 and 8×8 resolutions with increasing dilation rates [Chen et al., 2017] allowing exponentially increasing receptive fields.

We train these models for 500K steps, and a batch size of 512. Table 5 reports the performance of different configurations for inpainting. Global self-attention offers much better performance than the alternatives, re-affirming the importance of self-attention layers for such architectures. Surprisingly, local self-attention performs worse than fully-convolutional alternatives. We leave more detailed analysis of local self-attention in image generation architectures to future work.

Architecture	Model	FID ↓	IS ↑	PD ↓
<i>Fully Convolutional</i>	Dilated Convolutions	8.0	157.5	70.6
	More ResNet Blocks	8.1	157.1	71.9
<i>Self-Attention</i>	Local Self-Attention	9.4	149.8	78.2
	Global Self-Attention	7.4	164.8	67.1

Table 5: Fully Convolutional vs. Self-Attention architectures for inpainting.

5.6 Sample Diversity

In this section, we analyze the sample diversity of Palette on two translation tasks, namely, colorization and inpainting. We analyze the impact of the changing the diffusion objective function L_{simple} [Ho



Figure 7: Diversity of Palette outputs on colorization (top), inpainting (middle) and uncropping (bottom). Figures C.3, C.5, C.8 and C.9 have more samples.

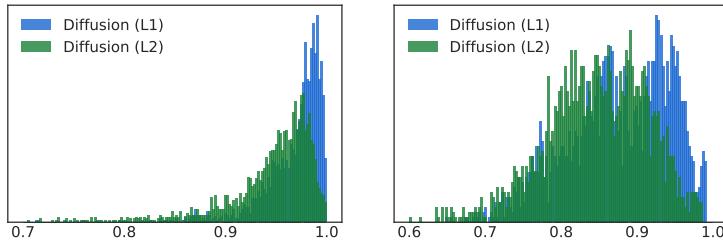


Figure 8: Distribution of pairwise multi-scale SSIM for Colorization (left) and Inpainting (right). Figure

et al., 2020], and compare L_1 vs. L_2 on sample diversity. While some existing conditional generative diffusion models such as SR3 [Saharia et al., 2021a] and WaveGrad [Chen et al., 2021a] have found L_1 norm to perform better than the originally proposed L_2 objective, there has not been a detailed analysis of the impact of these two objectives.

Inspired by Guadarrama et al. [2017], we generate multiple outputs for a given input, and then compute pairwise multi-scale SSIM between the first output sample and the remaining samples. We do this for multiple inputs, and then plot the histogram of SSIM values in Figure 8, for colorization and inpainting. Higher SSIM scores imply Less diversity. One can see that models trained with the L_2 loss have lower SSIM scores among their output samples than those trained with the L_1 loss, indicating greater sample diversity.

5.7 Multi-Task Learning

Multi-task training for image-to-image translation is a challenging but largely unexplored area of research. In prior work, the training paradigm is often restricted to multiple levels of corruptions within a single translation task – e.g., Yu et al. [2019, 2018a] train inpainting models on different amounts of masked regions, and Galteri et al. [2019], Liu et al. [2018] train JPEG artifact removal models on multiple quality factors. In this section, we train a single generalist Palette model on multiple translation tasks simultaneously. Specifically, we train Palette on tasks of JPEG decompression, inpainting, and colorization.

Table 6 indicates that multi-task Palette outperforms the task-specific model on JPEG decompression, it lags behind task-specific Palette models on inpainting and colorization. Due to compute limitations, we currently train multi-task Palette for the same number of training steps as task-specific models. We expect its performance to improve with more training.

Task	Model	FID ↓	IS ↑	CA ↑	PD ↓
<i>Inpainting</i> <i>(128×128 center mask)</i>	Palette (<i>Task-specific</i>)	6.6	173.9	69.3%	59.5
	Palette (<i>Multi-task</i>)	6.8	165.7	68.9%	65.2
<i>Colorization</i>	Regression (<i>Task-specific</i>)	5.5	176.9	68.0%	61.1
	Palette (<i>Task-specific</i>)	3.4	212.9	72.0%	48.0
	Palette (<i>Multi-task</i>)	3.7	187.4	69.4%	57.1
<i>JPEG decompression</i> <i>(QF = 5)</i>	Regression (<i>Task-specific</i>)	29.0	73.9	52.8%	155.4
	Palette (<i>Task-specific</i>)	8.3	133.6	64.2%	95.5
	Palette (<i>Multi-Task</i>)	7.0	137.8	64.7%	92.4

Table 6: Performance of Multi-Task Palette on various image-to-image translation tasks.

6 Conclusion

This paper shows that conditional diffusion models present a flexible and generic solution to image-to-image translation, rivaling the quality of GANs. Palette advances state-of-the-art on four challenging tasks and produces diverse outputs consistent with different input contexts. This is accomplished without task-specific customization or optimization instability. We also take a first step toward building multi-task image-to-image diffusion models that learn representations useful for a range of enhancement tasks. Further exploration and investigation of multi-task diffusion models is an exciting avenue for future work. Before Palette is ready for practical applications, its inherent biases should be studied and mitigated, its inference speed improved, and its operating resolution (currently 256×256) increased.

7 Acknowledgements

We thank Jon Shlens, Doug Eck, Rif A. Saurous, Hossein Talebi and the Google Brain team for research discussions and technical assistance.

References

- Eirikur Agustsson and Radu Timofte. NTIRE 2017 Challenge on Single Image Super-Resolution: Dataset and Study. In *CVPRW*, 2017.
- Lynton Ardizzone, Carsten Lüth, Jakob Kruse, Carsten Rother, and Ullrich Köthe. Guided Image Generation with Conditional Invertible Neural Networks. In *arXiv:1907.02392*, 2019.
- Martin Arjovsky, Soumith Chintala, and Léon Bottou. Wasserstein GAN. In *arXiv*, 2017.
- Jacob Austin, Daniel Johnson, Jonathan Ho, Danny Tarlow, and Rianne van den Berg. Structured denoising diffusion models in discrete state-spaces. *arXiv preprint arXiv:2107.03006*, 2021.
- Connelly Barnes, Eli Shechtman, Adam Finkelstein, and Dan B Goldman. PatchMatch: A randomized correspondence algorithm for structural image editing. *ACM Transactions on Graphics (Proc. SIGGRAPH)*, 28(3), August 2009.
- Marcelo Bertalmio, Guillermo Sapiro, Vincent Caselles, and Coloma Ballester. Image inpainting. In *Proceedings of the 27th annual conference on Computer graphics and interactive techniques*, pages 417–424, 2000.
- Richard Strong Bowen, Huiwen Chang, Charles Herrmann, Piotr Teterwak, Ce Liu, and Ramin Zabih. Oconet: Image extrapolation by object completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2307–2317, 2021.
- Ruojin Cai, Guandao Yang, Hadar Averbuch-Elor, Zekun Hao, Serge Belongie, Noah Snavely, and Bharath Hariharan. Learning Gradient Fields for Shape Generation. In *ECCV*, 2020.
- Liang-Chieh Chen, George Papandreou, Florian Schroff, and Hartwig Adam. Rethinking atrous convolution for semantic image segmentation. *arXiv preprint arXiv:1706.05587*, 2017.
- Nanxin Chen, Yu Zhang, Heiga Zen, Ron J. Weiss, Mohammad Norouzi, and William Chan. WaveGrad: Estimating Gradients for Waveform Generation. In *ICLR*, 2021a.

- Nanxin Chen, Yu Zhang, Heiga Zen, Ron J. Weiss, Mohammad Norouzi, Najim Dehak, and William Chan. WaveGrad 2: Iterative Refinement for Text-to-Speech Synthesis . In *INTERSPEECH*, 2021b.
- Yen-Chi Cheng, Chieh Hubert Lin, Hsin-Ying Lee, Jian Ren, Sergey Tulyakov, and Ming-Hsuan Yang. In&out: Diverse image outpainting via gan inversion. *arXiv preprint arXiv:2104.00675*, 2021.
- Yunjey Choi, Minje Choi, Munyoung Kim, Jung-Woo Ha, Sunghun Kim, and Jaegul Choo. Stargan: Unified generative adversarial networks for multi-domain image-to-image translation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 8789–8797, 2018.
- Ryan Dahl, Mohammad Norouzi, and Jonathon Shlens. Pixel recursive super resolution. In *ICCV*, 2017.
- Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE conference on computer vision and pattern recognition*, pages 248–255. Ieee, 2009.
- Aditya Deshpande, Jiajun Lu, Mao-Chuang Yeh, Min Jin Chong, and David Forsyth. Learning diverse image colorization. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 6837–6845, 2017.
- Prafulla Dhariwal and Alex Nichol. Diffusion models beat gans on image synthesis. *arXiv preprint arXiv:2105.05233*, 2021.
- Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. Density estimation using real NVP. *arXiv:1605.08803*, 2016.
- Chao Dong, Yubin Deng, Chen Change Loy, and Xiaoou Tang. Compression artifacts reduction by a deep convolutional network. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 576–584, 2015.
- Max Ehrlich, Larry Davis, Ser-Nam Lim, and Abhinav Shrivastava. Quantization guided jpeg artifact correction. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part VIII 16*, pages 293–309. Springer, 2020.
- Leonardo Galteri, Lorenzo Seidenari, Marco Bertini, and Alberto Del Bimbo. Deep generative adversarial compression artifact removal. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 4826–4835, 2017.
- Leonardo Galteri, Lorenzo Seidenari, Marco Bertini, and Alberto Del Bimbo. Deep universal generative adversarial compression artifact removal. *IEEE Transactions on Multimedia*, 21(8):2131–2145, 2019.
- Ian J Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative Adversarial Networks. *NIPS*, 2014.
- Sergio Guadarrama, Ryan Dahl, David Bieber, Mohammad Norouzi, Jonathon Shlens, and Kevin Murphy. Pixcolor: Pixel recursive colorization. *arXiv preprint arXiv:1705.07208*, 2017.
- Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, and Aaron Courville. Improved training of wasserstein gans. *arXiv preprint arXiv:1704.00028*, 2017.
- James Hays and Alexei A Efros. Scene completion using millions of photographs. *ACM Transactions on Graphics (ToG)*, 26(3):4–es, 2007.
- Kaiming He and Jian Sun. Statistics of patch offsets for image completion. In *European conference on computer vision*, pages 16–29. Springer, 2012.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *arXiv preprint arXiv:2006.11239*, 2020.
- Jonathan Ho, Chitwan Saharia, William Chan, David J. Fleet, Mohammad Norouzi, and Tim Salimans. Cascaded Diffusion Models for High Fidelity Image Generation. In *arXiv*, 2021.
- Emiel Hoogeboom, Didrik Nielsen, Priyank Jaini, Patrick Forré, and Max Welling. Argmax flows and multinomial diffusion: Towards non-autoregressive language models. *arXiv preprint arXiv:2102.05379*, 2021.
- Satoshi Iizuka, Edgar Simo-Serra, and Hiroshi Ishikawa. Globally and locally consistent image completion. *ACM Transactions on Graphics (ToG)*, 36(4):1–14, 2017.
- Phillip Isola, Jun-Yan Zhu, and Tinghui Zhou ajnd Alexei A. Efros. Image-to-Image Translation with Conditional Adversarial Nets. In *CVPR*, 2017a.

- Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A Efros. Image-to-image translation with conditional adversarial networks. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1125–1134, 2017b.
- Zahra Kadkhodaie and Eero P Simoncelli. Solving linear inverse problems using the prior implicit in a denoiser. *arXiv preprint 2007.13640*, 2021.
- Jiwon Kim, Jung Kwon Lee, and Kyoung Mu Lee. Deeply-recursive convolutional network for image super-resolution. In *CVPR*, pages 1637–1645, 2016.
- Soo Ye Kim, Kfir Aberman, Nori Kanazawa, Rahul Garg, Neal Wadhwa, Huiwen Chang, Nikhil Karnad, Munchurl Kim, and Orly Liba. Zoom-to-inpaint: Image inpainting with high-frequency details, 2021.
- Diederik P. Kingma and Prafulla Dhariwal. Glow: Generative Flow with Invertible 1x1 Convolutions. In *NIPS*, 2018.
- Diederik P Kingma and Max Welling. Auto-Encoding Variational Bayes. In *ICLR*, 2013.
- Diederik P Kingma, Tim Salimans, Ben Poole, and Jonathan Ho. Variational diffusion models. *arXiv preprint arXiv:2107.00630*, 2021.
- Zhifeng Kong, Wei Ping, Jiaji Huang, Kexin Zhao, and Bryan Catanzaro. Diffwave: A versatile diffusion model for audio synthesis. *arXiv preprint arXiv:2009.09761*, 2020.
- Zhifeng Kong, Wei Ping, Jiaji Huang, Kexin Zhao, and Bryan Catanzaro. DiffWave: A Versatile Diffusion Model for Audio Synthesis. *ICLR*, 2021.
- Johannes Kopf, Wolf Kienzle, Steven Drucker, and Sing Bing Kang. Quality prediction for image completion. *ACM Transactions on Graphics (ToG)*, 31(6):1–8, 2012.
- Manoj Kumar, Dirk Weissenborn, and Nal Kalchbrenner. Colorization transformer. In *ICLR 2021*, 2021. URL <https://openreview.net/forum?id=5NA1PinlGFu>.
- Gustav Larsson, Michael Maire, and Gregory Shakhnarovich. Learning representations for automatic colorization. In *European conference on computer vision*, pages 577–593. Springer, 2016.
- Christian Ledig, Lucas Theis, Ferenc Huszár, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken, Alykhan Tejani, Johannes Totz, Zehan Wang, et al. Photo-realistic single image super-resolution using a generative adversarial network. In *ICCV*, 2017.
- Chieh Hubert Lin, Hsin-Ying Lee, Yen-Chi Cheng, Sergey Tulyakov, and Ming-Hsuan Yang. Infinitygan: Towards infinite-resolution image synthesis. *arXiv preprint arXiv:2104.03963*, 2021.
- Hongyu Liu, Bin Jiang, Yibing Song, Wei Huang, and Chao Yang. Rethinking image inpainting via a mutual encoder-decoder with feature equalizations. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part II 16*, pages 725–741. Springer, 2020.
- Ming-Yu Liu, Xun Huang, Arun Mallya, Tero Karras, Timo Aila, Jaakko Lehtinen, and Jan Kautz. Few-shot unsupervised image-to-image translation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 10551–10560, 2019.
- Pengju Liu, Hongzhi Zhang, Kai Zhang, Liang Lin, and Wangmeng Zuo. Multi-level wavelet-cnn for image restoration. In *Proceedings of the IEEE conference on computer vision and pattern recognition workshops*, pages 773–782, 2018.
- D. Martin, C. Fowlkes, D. Tal, and J. Malik. A database of human segmented natural images and its application to evaluating segmentation algorithms and measuring ecological statistics. In *ICCV*, 2001.
- Chenlin Meng, Yang Song, Jiaming Song, Jiajun Wu, Jun-Yan Zhu, and Stefano Ermon. Sdedit: Image synthesis and editing with stochastic differential equations. *arXiv preprint arXiv:2108.01073*, 2021.
- Sachit Menon, Alexandru Damian, Shijia Hu, Nikhil Ravi, and Cynthia Rudin. PULSE: Self-supervised photo upsampling via latent space exploration of generative models. In *CVPR*, 2020.
- Luke Metz, Ben Poole, David Pfau, and Jascha Sohl-Dickstein. Unrolled generative adversarial networks. *arXiv preprint arXiv:1611.02163*, 2016.
- Kamyar Nazeri, Eric Ng, Tony Joseph, Faisal Qureshi, and Mehran Ebrahimi. Edgeconnect: Structure guided image inpainting using edge prediction. In *Proceedings of the IEEE/CVF International Conference on Computer Vision Workshops*, pages 0–0, 2019.

- Alex Nichol and Prafulla Dhariwal. Improved denoising diffusion probabilistic models. *arXiv preprint arXiv:2102.09672*, 2021.
- Niki Parmar, Ashish Vaswani, Jakob Uszkoreit, Lukasz Kaiser, Noam Shazeer, Alexander Ku, and Dustin Tran. Image transformer. In *ICML*, 2018.
- Guocheng Qian, Jinjin Gu, Jimmy Ren, Chao Dong, Furong Zhao, and Juan Lin. Trinity of Pixel Enhancement: a Joint Solution for Demosaicking, Denoising and Super-Resolution. In *arXiv:1905.02538*, 2019.
- Alec Radford, Luke Metz, and Soumith Chintala. Unsupervised representation learning with deep convolutional generative adversarial networks. *arXiv preprint arXiv:1511.06434*, 2015.
- Suman Ravuri and Oriol Vinyals. Classification accuracy score for conditional generative models. *arXiv preprint arXiv:1905.10887*, 2019.
- Stefan Roth and Michael J Black. Fields of experts: A framework for learning image priors. In *2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'05)*, volume 2, pages 860–867. IEEE, 2005.
- Amelie Royer, Alexander Kolesnikov, and Christoph H. Lampert. Probabilistic Image Colorization. In *arXiv:1705.04258*, 2017.
- Chitwan Saharia, Jonathan Ho, William Chan, Tim Salimans, David J. Fleet, and Mohammad Norouzi. Image Super-Resolution via Iterative Refinement. *arXiv preprint arXiv:2104.07636*, 2021a.
- Chitwan Saharia, Jonathan Ho, William Chan, Tim Salimans, David J Fleet, and Mohammad Norouzi. Image super-resolution via iterative refinement. *arXiv preprint arXiv:2104.07636*, 2021b.
- Tim Salimans, Andrej Karpathy, Xi Chen, and Diederik P. Kingma. PixelCNN++: Improving the PixelCNN with Discretized Logistic Mixture Likelihood and Other Modifications. In *ICLR*, 2017.
- Hiroshi Sasaki, Chris G Willcocks, and Toby P Breckon. Unit-ddpm: Unpaired image translation with denoising diffusion probabilistic models. *arXiv preprint arXiv:2104.05358*, 2021.
- Qi Shan, Brian Curless, Yasutaka Furukawa, Carlos Hernandez, and Steven M Seitz. Photo uncrop. In *European Conference on Computer Vision*, pages 16–31. Springer, 2014.
- Abhishek Sinha, Jiaming Song, Chenlin Meng, and Stefano Ermon. D2c: Diffusion-denoising models for few-shot conditional generation. *arXiv preprint arXiv:2106.06819*, 2021.
- Jascha Sohl-Dickstein, Eric Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. In *ICML*, pages 2256–2265. PMLR, 2015.
- Yang Song and Stefano Ermon. Improved Techniques for Training Score-Based Generative Models. *arXiv preprint arXiv:2006.09011*, 2020.
- Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. *arXiv preprint arXiv:2011.13456*, 2020.
- Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-Based Generative Modeling through Stochastic Differential Equations. In *ICLR*, 2021.
- Yaniv Taigman, Adam Polyak, and Lior Wolf. Unsupervised cross-domain image generation. *arXiv preprint arXiv:1611.02200*, 2016.
- Piotr Teterwak, Aaron Sarna, Dilip Krishnan, Aaron Maschinot, David Belanger, Ce Liu, and William T Freeman. Boundless: Generative adversarial networks for image extension. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 10521–10530, 2019.
- Arash Vahdat and Jan Kautz. NVAE: A deep hierarchical variational autoencoder. In *NeurIPS*, 2020.
- Arash Vahdat, Karsten Kreis, and Jan Kautz. Score-based generative modeling in latent space. *arXiv preprint arXiv:2106.05931*, 2021.
- Aaron van den Oord, Nal Kalchbrenner, Oriol Vinyals, Lasse Espeholt, Alex Graves, and Koray Kavukcuoglu. Conditional image generation with PixelCNN decoders. In *NIPS*, pages 4790–4798, 2016.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention Is All You Need. In *NIPS*, 2017.

- Ashish Vaswani, Prajit Ramachandran, Aravind Srinivas, Niki Parmar, Blake Hechtman, and Jonathon Shlens. Scaling local self-attention for parameter efficient visual backbones. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 12894–12904, 2021.
- Miao Wang, Yu-Kun Lai, Yuan Liang, Ralph R Martin, and Shi-Min Hu. Biggerpicture: data-driven image extrapolation using graph matching. *ACM Transactions on Graphics*, 33(6), 2014.
- Yi Wang, Xin Tao, Xiaoyong Shen, and Jiaya Jia. Wide-context semantic image extrapolation. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 1399–1408, 2019a.
- Yi Wang, Xin Tao, Xiaoyong Shen, and Jiaya Jia. Wide-context semantic image extrapolation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1399–1408, 2019b.
- Zongxin Yang, Jian Dong, Ping Liu, Yi Yang, and Shuicheng Yan. Very long natural scenery image prediction by outpainting. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 10561–10570, 2019.
- Zili Yi, Qiang Tang, Shekoofeh Azizi, Daesik Jang, and Zhan Xu. Contextual residual aggregation for ultra high-resolution image inpainting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7508–7517, 2020.
- Jiahui Yu, Zhe Lin, Jimei Yang, Xiaohui Shen, Xin Lu, and Thomas S Huang. Generative image inpainting with contextual attention. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 5505–5514, 2018a.
- Jiahui Yu, Zhe Lin, Jimei Yang, Xiaohui Shen, Xin Lu, and Thomas S Huang. Free-form image inpainting with gated convolution. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 4471–4480, 2019.
- Ke Yu, Chao Dong, Liang Lin, and Chen Change Loy. Crafting a toolchain for image restoration by deep reinforcement learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 2443–2452, 2018b.
- Richard Zhang, Phillip Isola, and Alexei A Efros. Colorful image colorization. In *European conference on computer vision*, pages 649–666. Springer, 2016.
- Shengyu Zhao, Jonathan Cui, Yilun Sheng, Yue Dong, Xiao Liang, Eric I Chang, and Yan Xu. Large scale image completion via co-modulated generative adversarial networks. *arXiv preprint arXiv:2103.10428*, 2021.
- Chuanxia Zheng, Tat-Jen Cham, and Jianfei Cai. Pluralistic image completion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 1438–1447, 2019.
- Bolei Zhou, Agata Lapedriza, Aditya Khosla, Aude Oliva, and Antonio Torralba. Places: A 10 million image database for scene recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2017.
- Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision*, pages 2223–2232, 2017.

A Diffusion Models

Diffusion models comprise a forward diffusion process and a reverse denoising process that is used at generation time. The forward diffusion process is a Markovian process that iteratively adds Gaussian noise to a data point $\mathbf{y}_0 \equiv \mathbf{y}$ over T iterations:

$$q(\mathbf{y}_{t+1} | \mathbf{y}_t) = \mathcal{N}(\mathbf{y}_{t-1}; \sqrt{\alpha_t} \mathbf{y}_{t-1}, (1 - \alpha_t)I) \quad (2)$$

$$q(\mathbf{y}_{1:T} | \mathbf{y}_0) = \prod_{t=1}^T q(\mathbf{y}_t | \mathbf{y}_{t-1}) \quad (3)$$

where α_t are hyper-parameters of the noise schedule. The forward process with α_t is constructed in a manner where at $t = T$, \mathbf{y}_T is virtually indistinguishable from Gaussian noise. Note, we can also marginalize the forward process at each step:

$$q(\mathbf{y}_t | \mathbf{y}_0) = \mathcal{N}(\mathbf{y}_t; \sqrt{\gamma_t} \mathbf{y}_0, (1 - \gamma_t)I), \quad (4)$$

where $\gamma_t = \prod_{t'}^t \alpha'_{t'}$.

The Gaussian parameterization of the forward process also allows a closed form formulation of the posterior distribution of \mathbf{y}_{t-1} given $(\mathbf{y}_0, \mathbf{y}_t)$ as

$$q(\mathbf{y}_{t-1} | \mathbf{y}_0, \mathbf{y}_t) = \mathcal{N}(\mathbf{y}_{t-1} | \boldsymbol{\mu}, \sigma^2 \mathbf{I}) \quad (5)$$

where $\boldsymbol{\mu} = \frac{\sqrt{\gamma_{t-1}}(1-\alpha_t)}{1-\gamma_t} \mathbf{y}_0 + \frac{\sqrt{\alpha_t}(1-\gamma_{t-1})}{1-\gamma_t} \mathbf{y}_t$ and $\sigma^2 = \frac{(1-\gamma_{t-1})(1-\alpha_t)}{1-\gamma_t}$. This result proves to be very helpful during inference as shown below.

Learning: Palette learns a reverse process which inverts the forward process. Given a noisy image $\tilde{\mathbf{y}}$,

$$\tilde{\mathbf{y}} = \sqrt{\gamma} \mathbf{y}_0 + \sqrt{1-\gamma} \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad (6)$$

the goal is to recover the target image \mathbf{y}_0 . We parameterize our neural network model $f_\theta(x, \tilde{\mathbf{y}}, \gamma)$ to condition on the input x , a noisy image $\tilde{\mathbf{y}}$, and the current noise level γ . Learning entails prediction of the noise vector $\boldsymbol{\epsilon}$ by optimizing the objective

$$\mathbb{E}_{(\mathbf{x}, \mathbf{y})} \mathbb{E}_{\boldsymbol{\epsilon}, \gamma} \left\| f_\theta(\mathbf{x}, \underbrace{\sqrt{\gamma} \mathbf{y}_0 + \sqrt{1-\gamma} \boldsymbol{\epsilon}}_{\tilde{\mathbf{y}}}, \gamma) - \boldsymbol{\epsilon} \right\|_p^p. \quad (7)$$

This objective, also known as L_{simple} in Ho et al. [2020], is equivalent to maximizing a weighted variational lower-bound on the likelihood [Ho et al., 2020].

Inference: Palette performs inference via the learned reverse process. Since the forward process is constructed so the prior distribution $p(\mathbf{y}_T)$ approximates a standard normal distribution $\mathcal{N}(\mathbf{y}_T | \mathbf{0}, \mathbf{I})$, the sampling process can start at pure Gaussian noise, followed by T steps of iterative refinement.

Also recall that the neural network model f_θ is trained to estimate $\boldsymbol{\epsilon}$, given any noisy image $\tilde{\mathbf{y}}$, and \mathbf{y}_t . Thus, given \mathbf{y}_t , we approximate \mathbf{y}_0 by rearranging terms in equation 6 as

$$\hat{\mathbf{y}}_0 = \frac{1}{\sqrt{\gamma_t}} \left(\mathbf{y}_t - \sqrt{1-\gamma_t} f_\theta(\mathbf{x}, \mathbf{y}_t, \gamma_t) \right). \quad (8)$$

Following Ho et al. [2020], we substitute our estimate $\hat{\mathbf{y}}_0$ into the posterior distribution of $q(\mathbf{y}_{t-1} | \mathbf{y}_0, \mathbf{y}_t)$ in equation 5 to parameterize the mean of $p_\theta(\mathbf{y}_{t-1} | \mathbf{y}_t, \mathbf{x})$ as

$$\mu_\theta(\mathbf{x}, \mathbf{y}_t, \gamma_t) = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{y}_t - \frac{1 - \alpha_t}{\sqrt{1 - \gamma_t}} f_\theta(\mathbf{x}, \mathbf{y}_t, \gamma_t) \right). \quad (9)$$

And we set the variance of $p_\theta(\mathbf{y}_{t-1} | \mathbf{y}_t, \mathbf{x})$ to $(1 - \alpha_t)$, a default given by the variance of the forward process Ho et al. [2020].

With this parameterization, each iteration of the reverse process can be computed as

$$\mathbf{y}_{t-1} \leftarrow \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{y}_t - \frac{1 - \alpha_t}{\sqrt{1 - \gamma_t}} f_\theta(\mathbf{x}, \mathbf{y}_t, \gamma_t) \right) + \sqrt{1 - \alpha_t} \boldsymbol{\epsilon}_t,$$

where $\boldsymbol{\epsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. This resembles one step of Langevin dynamics for which f_θ provides an estimate of the gradient of the data log-density.

Algorithm 1 Training a denoising model f_θ

```

1: repeat
2:    $(\mathbf{x}, \mathbf{y}_0) \sim p(\mathbf{x}, \mathbf{y})$ 
3:    $\gamma \sim p(\gamma)$ 
4:    $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
5:   Take a gradient descent step on
         $\nabla_\theta \|f_\theta(\mathbf{x}, \sqrt{\gamma}\mathbf{y}_0 + \sqrt{1-\gamma}\epsilon, \gamma) - \epsilon\|_p^p$ 
6: until converged

```

Algorithm 2 Inference in T iterative refinement steps

```

1:  $\mathbf{y}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{y}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{y}_t - \frac{1-\alpha_t}{\sqrt{1-\gamma_t}} f_\theta(\mathbf{x}, \mathbf{y}_t, \gamma_t) \right) + \sqrt{1-\alpha_t} \mathbf{z}$ 
5: end for
6: return  $\mathbf{y}_0$ 

```

B Implementation Details

Training Details : We train all models with a mini batch-size of 1024 for 1M training steps. We do not find over fitting to be an issue, and hence use the model checkpoint at 1M steps for reporting the final results. Consistent with previous works [Ho et al., 2020, Saharia et al., 2021a], we use standard Adam optimizer with a fixed 1e-4 learning rate and 10k linear learning rate warmup schedule. We use 0.9999 EMA for all our experiments. We do not perform any task-specific hyper-parameter tuning, or architectural modifications.

Diffusion Hyper-parameters : Following [Saharia et al., 2021a, Chen et al., 2021a] we use α conditioning for training Palette. This allows us to perform hyper-parameter tuning over noise schedules and refinement steps for Palette during inference. During training, we use a linear noise schedule of $(1e^{-6}, 0.01)$ with 2000 time-steps, and use 1000 refinement steps with a linear schedule of $(1e^{-4}, 0.09)$ during inference.

Task Specific Details:

- **Colorization** : We use RGB parameterization for colorization. We use the grayscale image as the source image and train Palette to predict the RGB image.
- **Inpainting** : We train Palette on a combination of free-form and rectangular masks. For free-form masks, we use Algorithm 1 in Yu et al. [2019]. For rectangular masks, we uniformly sample between 1 and 5 masks. The total area covered by the rectangular masks is kept between a total of 10% to 30% of the image. We randomly sample a free-form mask with 60% probability, and rectangular masks with 40% probability. We do not provide any additional mask channel, and simply fill the masked region with random Gaussian noise. During training, we mask restrict the L_{simple} loss function to the spatial region corresponding to masked regions, and use the model’s prediction for only the masked region during inference.
- **Uncropping** : We train the model for image extension along all four directions, or just one direction. In both cases, we set the masked region to 50% of the image. During training, we uniformly choose masking along one side, or masking along all 4 sides. When masking along one side, we further make a uniform random choice over the side. Rest of the training details are identical to inpainting.
- **JPEG Decompression** : We train Palette for JPEG Decompression on quality factors in (5, 30). Since decompression for lower quality factors is a significantly more difficult task, we use an exponential distribution to sample the quality factor during training. Specifically, the sampling probability of a quality range Q is set to $\propto e^{-\frac{Q}{10}}$.

C Additional Experimental Results

C.1 Colorization

Following prior work [Zhang et al., 2016, Guadarrama et al., 2017, Kumar et al., 2021], we train and evaluate models on ImageNet [Deng et al., 2009]. In order to compare our models with existing works in Table 1, we follow ColTran [Kumar et al., 2021] and use the first 5000 images from ImageNet validation set to report performance on standard metrics. We use the next 5000 images as the reference distribution for FID to mirror ColTran’s implementation. For benchmarking purposes, we also report the performance of Palette on ImageNet ctet10k [Larsson et al., 2016] dataset in Table C.1. Notice that while Palette trained with the L_2 loss produces more diverse samples than the L_1 loss, it lags behind slightly in terms of samples quality.

Human Evaluation: The ultimate evaluation of image-to-image translation models is human evaluation; *i.e.*, whether or not humans can discriminate model outputs from reference images. To this end we use controlled human experiments. In a series of two alternative forced choice trials, we ask subjects which of two side-by-side images is the real photo and which has been generated by the model. In particular, subjects are asked “*Which image would you guess is from a camera?*” Subjects viewed images for either 3 or 5 seconds before having to respond. For the experiments we compare outputs from four models against reference images, namely,

Model	FID-10K ↓	IS ↑	CA ↑	PD ↓
Palette (L_2)	3.4	212.9	72.0%	48.0
Palette (L_1)	3.4	215.8	71.9%	45.8
Ground Truth	2.7	250.1	76.0%	0.0

Table C.1: Benchmark numbers on ctest10k ImageNet subset for Image Colorization.

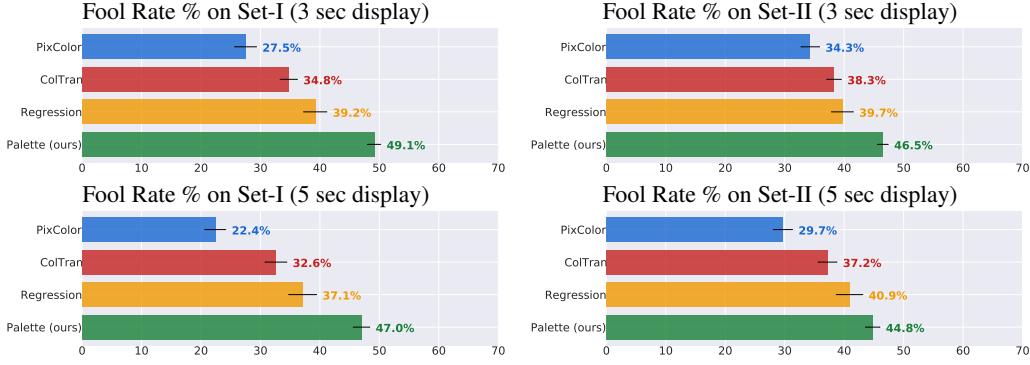


Figure C.1: Human evaluation results on ImageNet colorization.

PixColor [Guadarrama et al., 2017], Coltran [Kumar et al., 2021], our Regression baseline, and Palette. To summarize the result we compute the subject *fool rate*, i.e., the fraction of human raters who select the model outputs over the reference image. We use a total of 100 images for human evaluation, and divide these into two independent subsets - Set-I and Set-II, each of which is seen by 50 subjects.

As shown in Figure C.1, the fool rate for Palette is close to 50% and higher than baselines in all cases. We note that when subjects are given less time to inspect the images the fool rates are somewhat higher, as expected. We also note the strength of our regression baseline, which also performs better than PixColor and Coltran. Finally, to provide insight into the human evaluation results we also show several more examples of Palette output, with comparisons to benchmarks, in Figure C.2. One can see that in several cases, Palette has learned colors that are more meaningful and consistent with the reference images and the semantic content of the images. Figure C.3 also shows the natural diversity of Palette outputs for colorization model.

C.2 Inpainting

Comparison on 256×256 images: We report all inpainting results on 256×256 center cropped images. In order to ensure fair comparison, we use a fixed set of image-mask pair for each configuration for all models during evaluation. Since HiFill [Yi et al., 2020] is primarily trained on 512×512 images, we use 512×512 center crops with exact same mask within the central 256×256 region. This provides HiFill with 4× bigger inpainting context compared to DeepFillv2 and Palette.

We train two Palette models for Inpainting - i) Palette (I) trained on ImageNet dataset, and ii) Palette (I+P) trained on mixture of ImageNet and Places2 dataset. Table C.2 shows full comparison of Palette with existing methods on all inpainting configurations. Based on the type of mask, and the area covered, we report results for the following categories - i) 10-20% free-form region, ii) 20-30% free-form region, iii) 30-40% free-form region and iv) 128×128 center rectangle region. Palette consistently outperforms existing works by a significant margin on all configurations. Interestingly Palette (I) performs slightly better than Palette (I+P) on ImageNet indicating that augmentation with Places2 images during training doesn't boost ImageNet performance. Furthermore, Palette (I) is only slightly worse compared to Palette (I+P) on Places2 even though it is not trained on Places2 images. We observe a significant drop in the performance of HiFill [Yi et al., 2020] with larger masks. It is important to note that DeepFillv2 and HiFill are not trained on ImageNet, but we report their performance on ImageNet ctest10k primarily for benchmarking purposes.

C.3 Uncropping

Many existing uncropping methods [Cheng et al., 2021, Teterwak et al., 2019] have been trained on different subsets of Places2 [Zhou et al., 2017] dataset. In order to maintain uniformity, we follow a similar setup as inpainting and train Palette on a combined dataset of Places2 and ImageNet. While we train Palette to extend the image in all directions or just one direction, to compare fairly against existing methods we evaluate Palette

Mask Type	Model	ImageNet			Places2	
		FID ↓	IS ↑	CA ↑	PD ↓	PD ↓
10-20% <i>Free-Form Mask</i>	DeepFillv2 [Yu et al., 2019]	6.7	198.2	71.6%	38.6	12.2
	HiFill [Yi et al., 2020]	7.5	192.0	70.1%	46.9	13.0
	Palette (I) (Ours)	5.1	221.0	73.8%	15.6	11.6
	Palette (I+P) (Ours)	5.2	219.2	73.7%	15.5	11.6
20-30% <i>Free-Form Mask</i>	DeepFillv2 [Yu et al., 2019]	9.4	174.6	68.8%	64.7	13.5
	HiFill [Yi et al., 2020]	12.4	157.0	65.7%	86.2	15.7
	Palette (I) (Ours)	5.2	208.6	72.6%	27.4	11.8
	Palette (I+P) (Ours)	5.2	205.5	72.3%	27.6	11.7
30-40% <i>Free-Form Mask</i>	DeepFillv2 [Yu et al., 2019]	14.2	144.7	64.9%	95.5	15.8
	HiFill [Yi et al., 2020]	20.9	115.6	59.4%	131.0	20.1
	Palette (I)	5.5	195.2	71.4%	39.9	12.1
	Palette (I+P)	5.6	192.8	71.3%	40.2	11.6
<i>128×128 Center Mask</i>	DeepFillv2 [Yu et al., 2019]	18.0	135.3	64.3%	117.2	15.3
	HiFill [Yi et al., 2020]	20.1	126.8	62.3%	129.7	16.9
	Palette (I)	6.4	173.3	69.7%	58.8	12.2
	Palette (I+P)	6.6	173.9	69.3%	59.5	11.9
Ground Truth		5.1	231.6	74.6%	0.0	11.4
						0.0

Table C.2: Quantitative evaluation for inpainting on ImageNet and Places2 validation images.

on extending only the right half of the image. For Table 3, we use ctest10k and places10k to report results on ImageNet and Places2 validation sets respectively.

We also perform category specific evaluation of Palette with existing techniques - Boundless [Teterwak et al., 2019] and InfinityGAN [Lin et al., 2021]. Since Boundless is only trained on top-50 categories from Places2 dataset, we compare Palette with Boundless specifically on these categories from Places2 validation set in Table C.3. Palette achieves significantly better performance compared to Boundless re-affirming the strength of our model. Furthermore, we compare Palette with a more recent GAN based uncropping technique - InfinityGAN [Lin et al., 2021]. In order to fairly compare Palette with InfinityGAN, we specifically evaluate on the scenery categories from Places2 validation and test set. We use the samples generously provided by Lin et al. [2021], and generate outputs for Boundless, and Palette. Table C.4 shows that Palette is significantly better than domain specific model InfinityGAN on scenery images in terms of automated metrics.

Model	FID ↓	PD ↓
Boundless [Teterwak et al., 2019]	28.3	115.0
Palette	22.9	93.4
Ground Truth	23.6	0.0

Table C.3: Comparison with uncropping method Boundless [Teterwak et al., 2019] on top-50 Places2 categories.

Model	FID ↓
Boundless [Teterwak et al., 2019]	12.7
InfinityGAN [Lin et al., 2021]	15.7
Palette	5.6

Table C.4: Comparison with uncropping method InfinityGAN [Lin et al., 2021] and Boundless [Teterwak et al., 2019] on scenery categories.

Human Evaluation: Like colorization, we also report results from human evaluation experiments. Obtaining high fool rates for uncropping is a significantly more challenging task than colorization, because one half of the image area is fully generated by the model. As a consequence there are more opportunities for synthetic artifacts. Because the baselines available for uncropping are trained and tested on Places2, we run human evaluation experiments only on Places2. Beyond the choice of dataset, all other aspects of experimental design are identical to that used above for colorization, with two disjoint sets of test images, namely, Set-I and Set-II.

The results are characterized in terms of the fool rate, and are shown in Figure C.6. Palette obtains significantly higher fool rates on all human evaluation runs compared to existing methods, i.e., Boundless [Teterwak et al., 2019] and InfinityGAN [Lin et al., 2021]. Interestingly, when raters are given more time to inspect each pair of images, the fool rates for InfinityGAN and Boundless worsen considerably. Palette, on the other hand, observes approximately similar fool rates.

C.4 JPEG Decompression

In order to be consistent with other tasks, we perform training and evaluation on ImageNet dataset. Note that this is unlike most prior work [Dong et al., 2015, Liu et al., 2018], which mainly use small datasets such as DIV2K [Agustsson and Timofte, 2017] and BSD500 [Martin et al., 2001] for training and evaluation. Recent works such as [Galteri et al., 2019] use a relatively larger MS-COCO dataset for training, however, to the best of our knowledge, we are the first to train and evaluate JPEG artifact removal on ImageNet. We compare Palette with a strong Regression baseline which uses an identical architecture. We report results on JPEG quality factor settings of 5, 10 and 20 in Table 4.



Figure C.2: Comparison of different methods for colorization on ImageNet validation images. Baselines:
[†][Guadarrama et al., 2017] and [‡][Kumar et al., 2021].

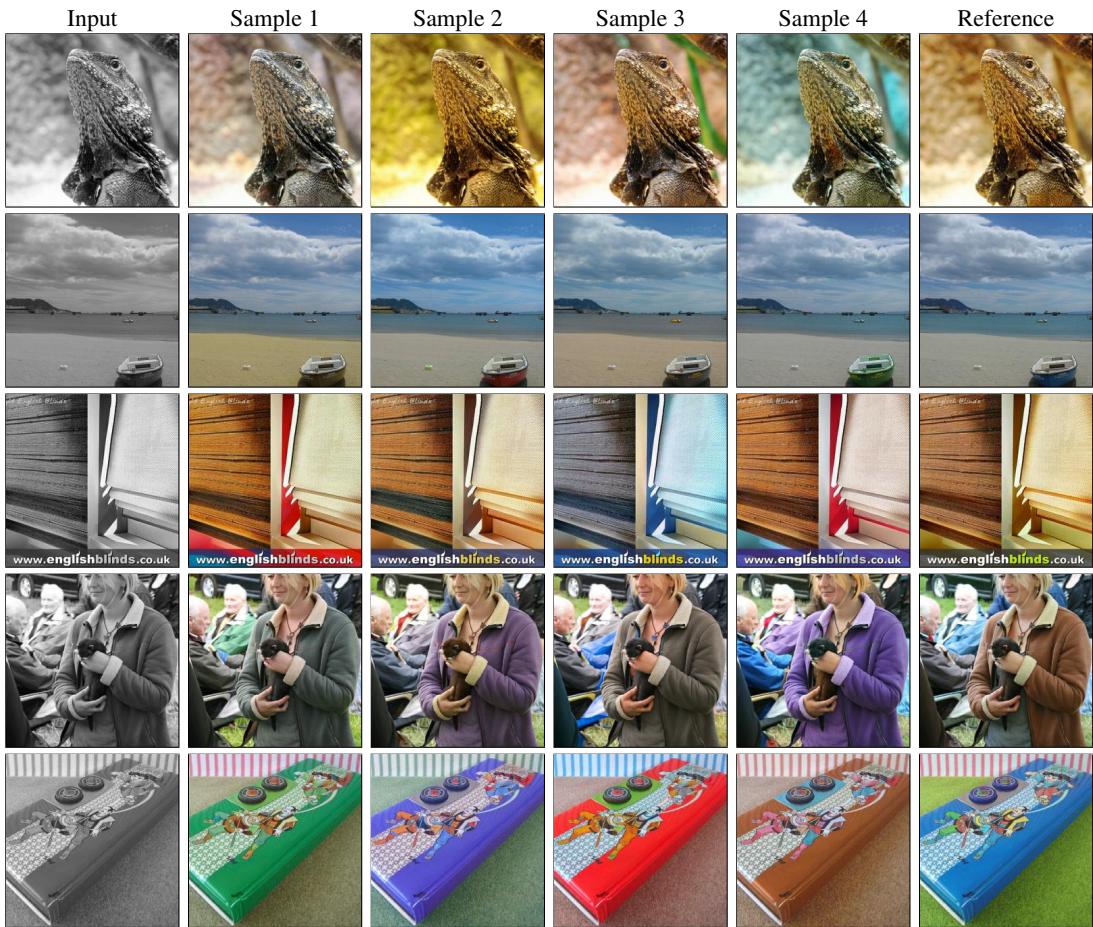


Figure C.3: Diversity of Palette outputs on ImageNet colorization validation images.

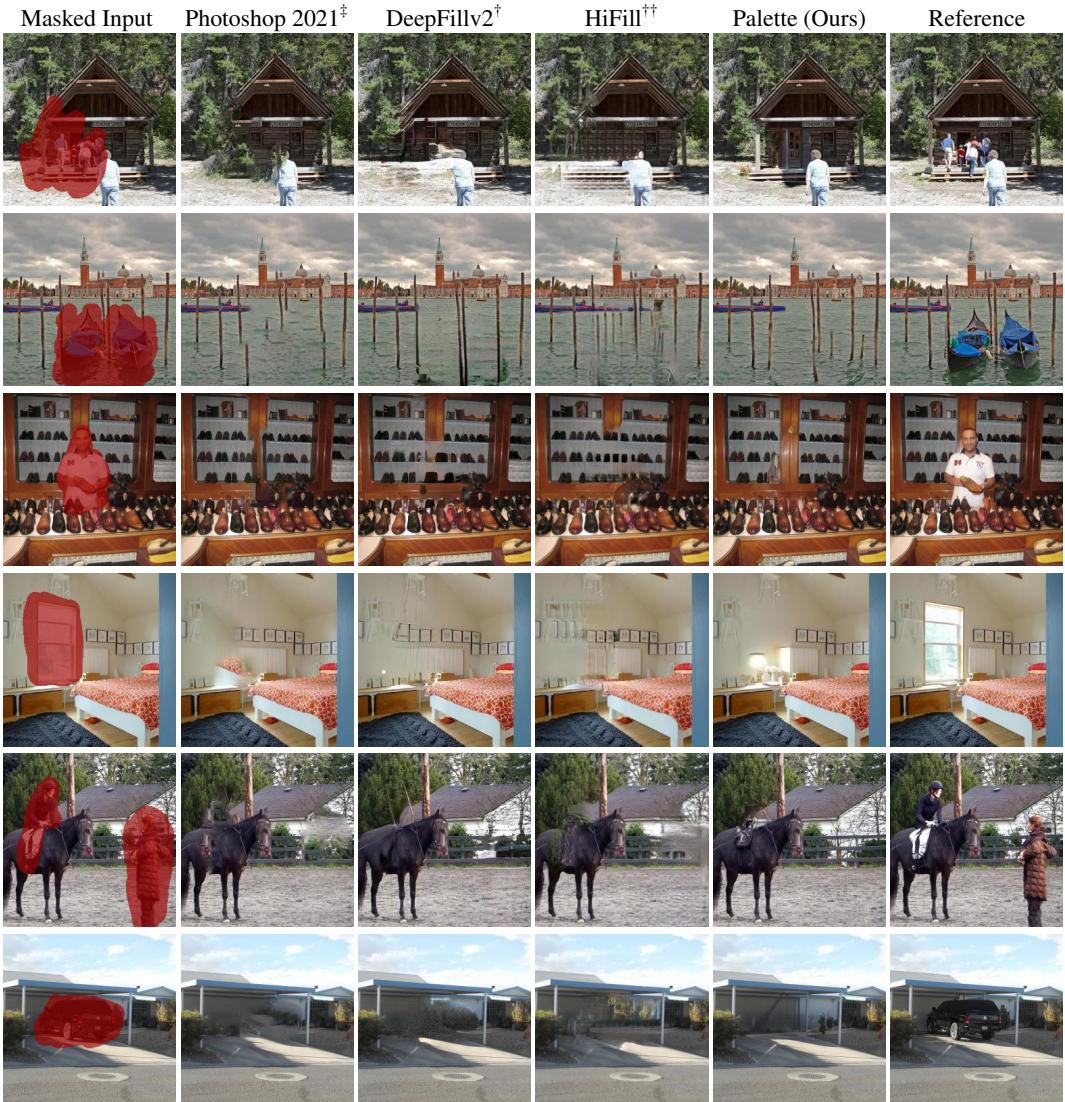


Figure C.4: Comparison of inpainting methods on object removal. Baselines: [‡]Photoshop’s *Content-aware Fill*, based on PatchMatch [Barnes et al., 2009], [†][Yu et al., 2019], and ^{††}[Yi et al., 2020].

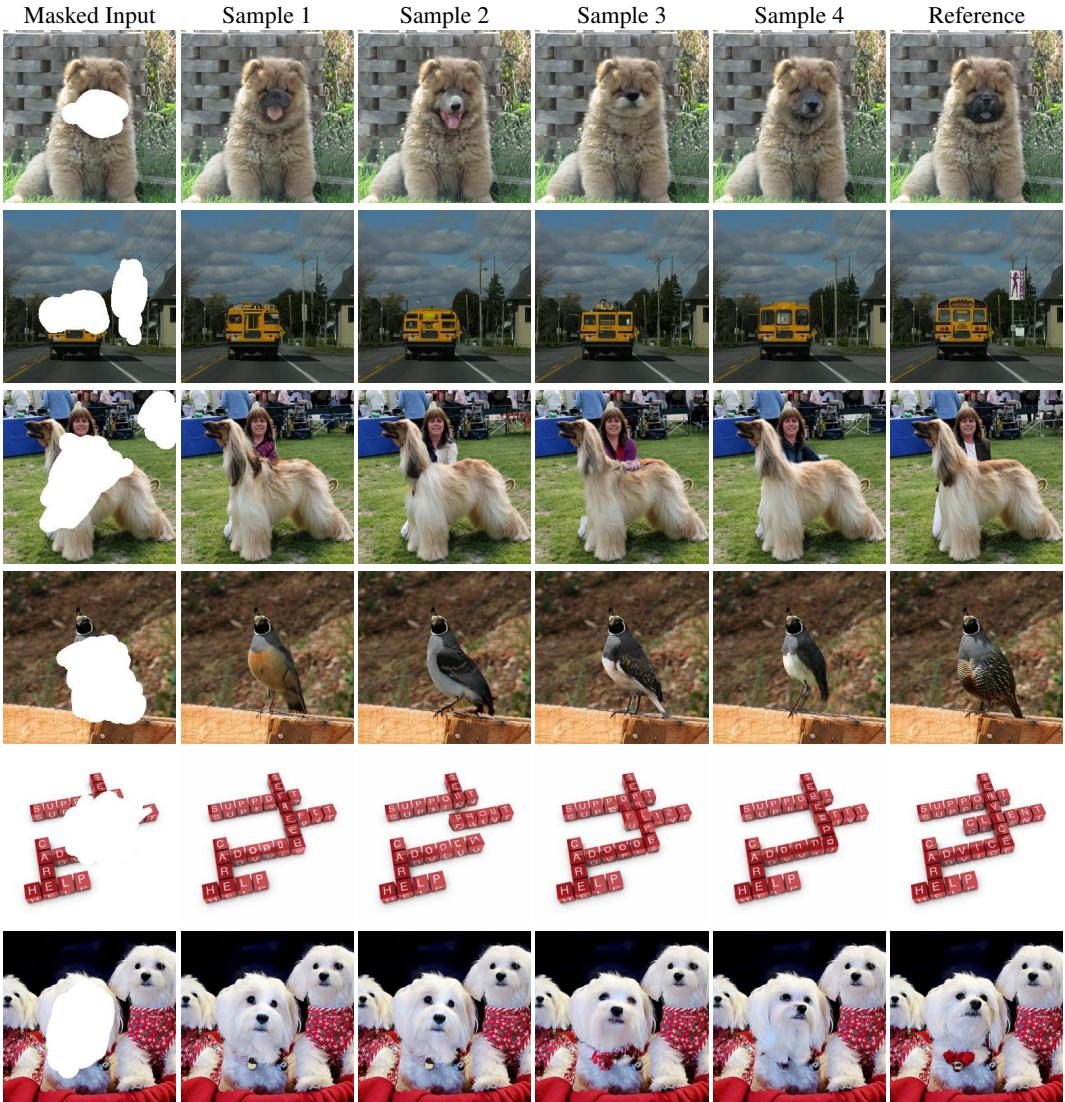


Figure C.5: Diversity of Palette outputs on image inpainting.

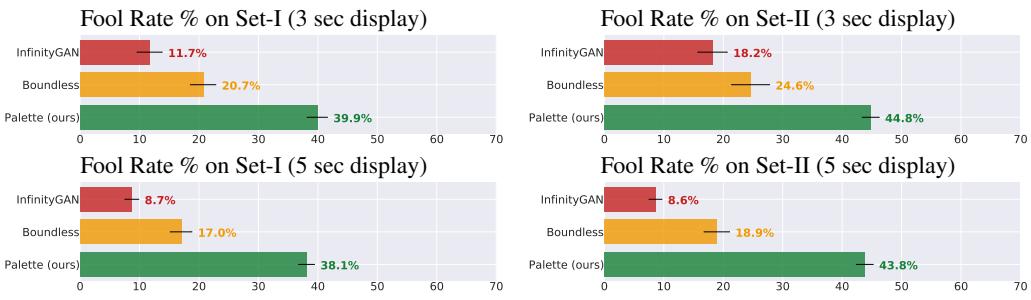


Figure C.6: Human evaluation results on Places2 uncropping.



Figure C.7: Image uncropping results on Places2 validation images. Baselines: Boundless[†] [Teterwak et al., 2019] and InfinityGAN^{††} [Lin et al., 2021] trained on a scenery subset of Places2. Samples for both baselines are generously provided by their respective authors.

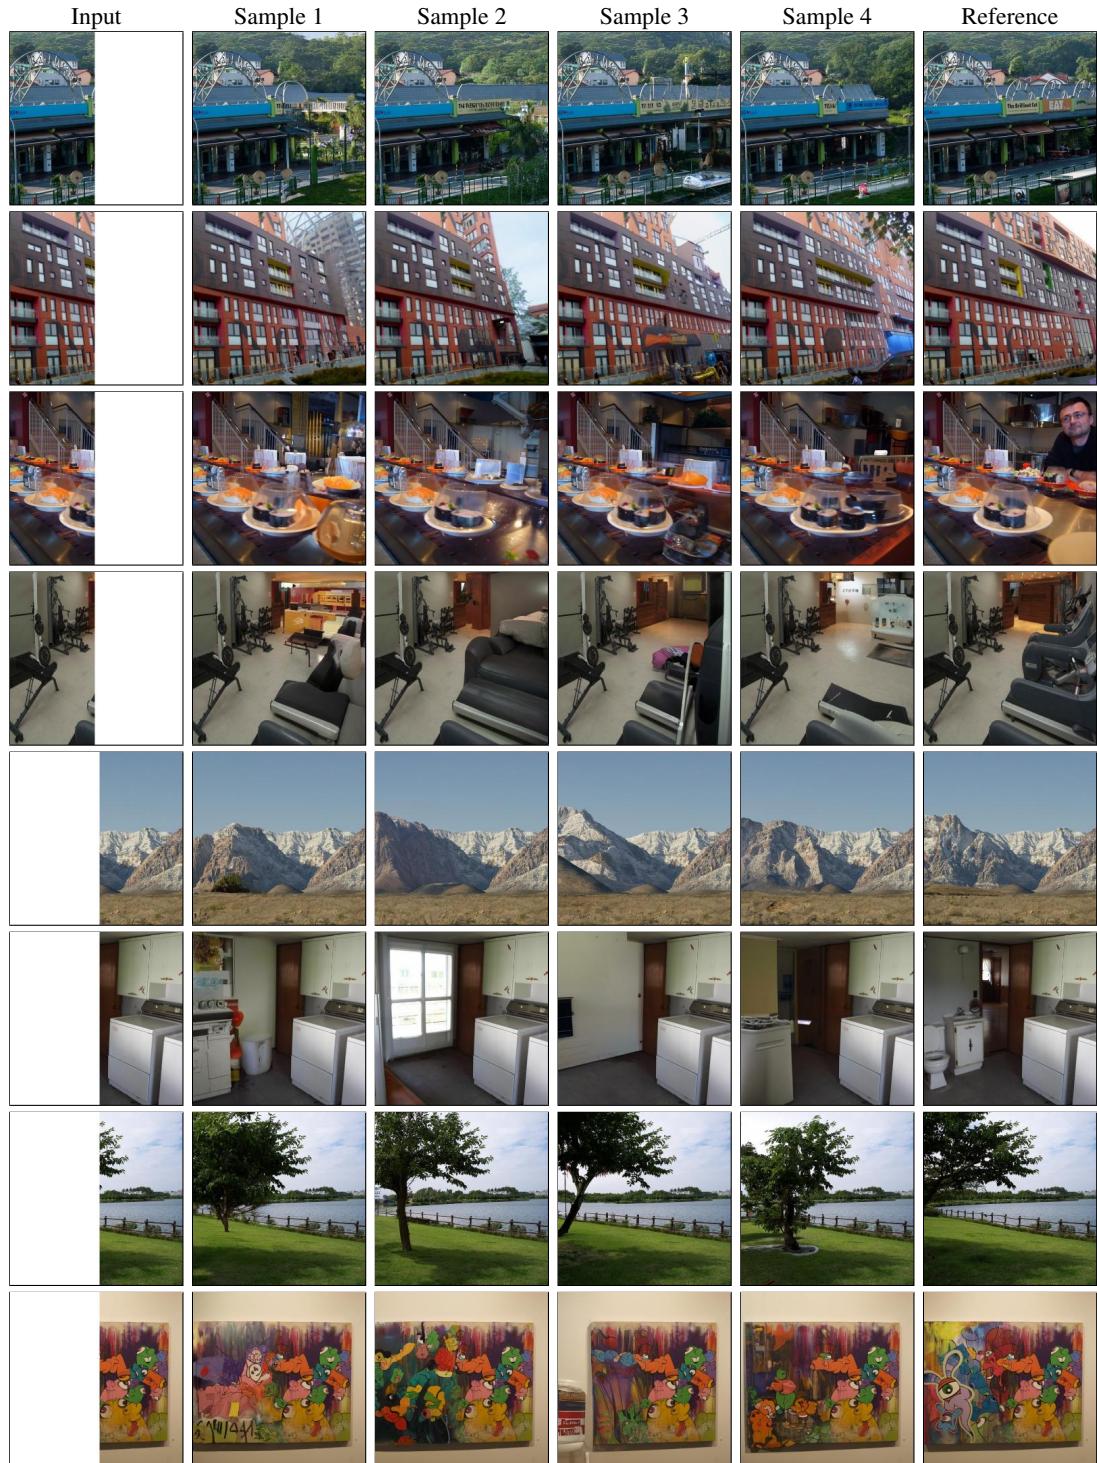


Figure C.8: Diversity of Palette outputs on Right and Left uncropping on Places2 dataset.

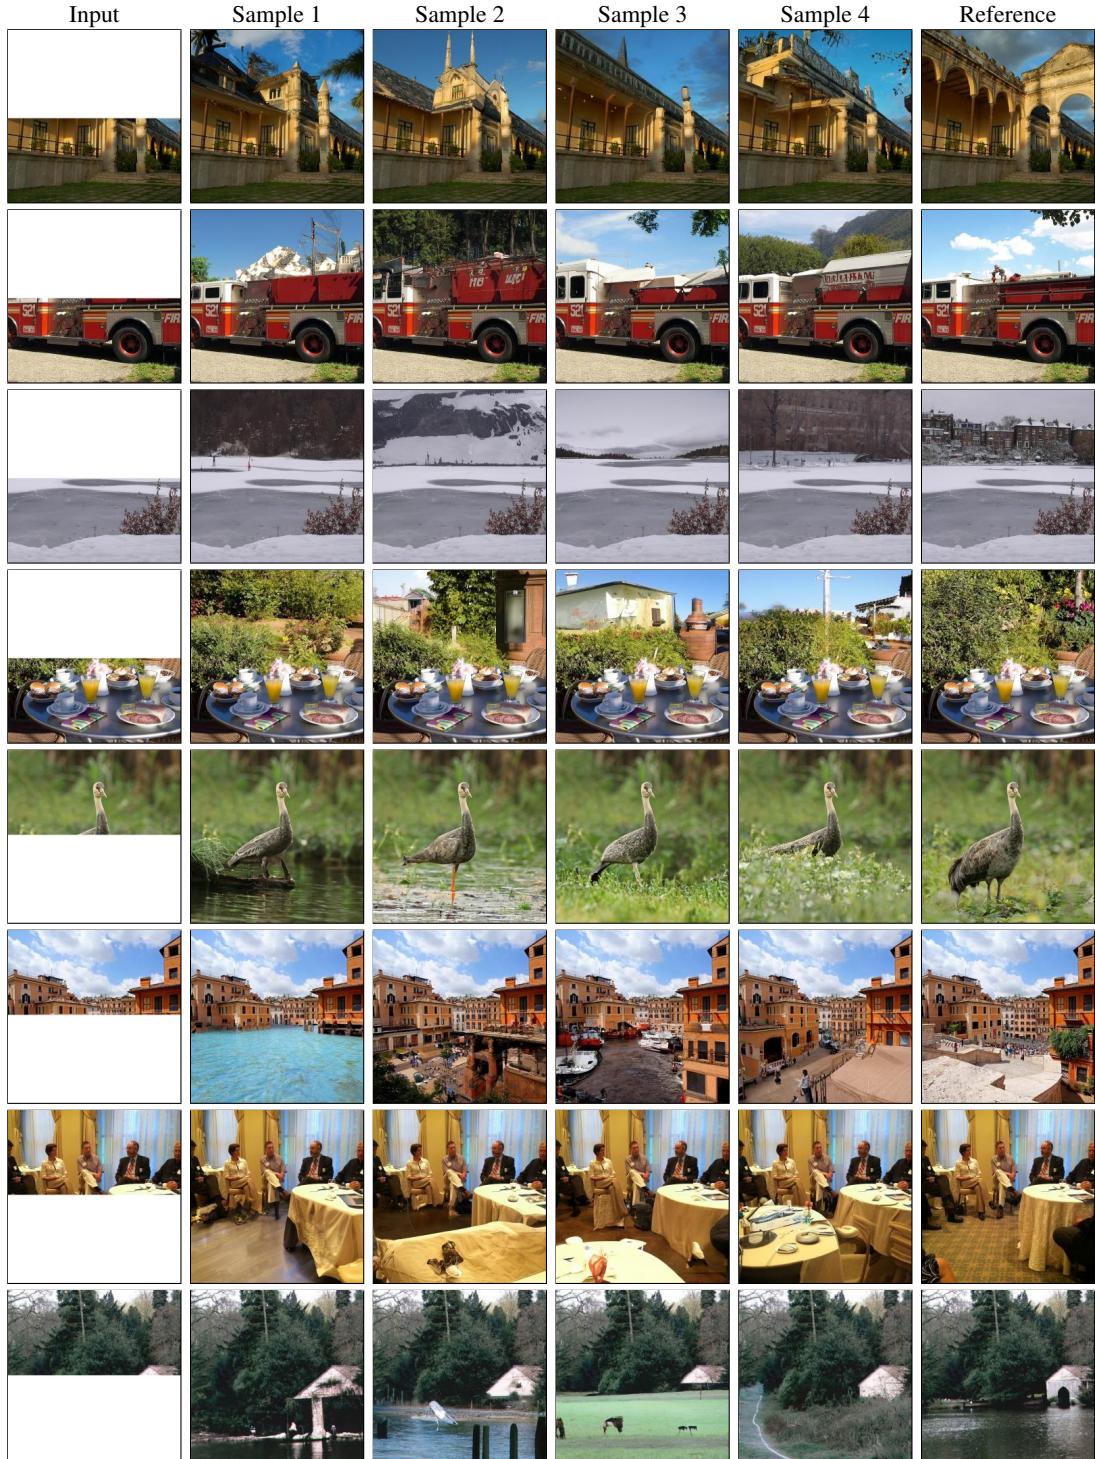


Figure C.9: Diversity of Palette outputs on Top and Bottom uncropping on Places2 dataset.

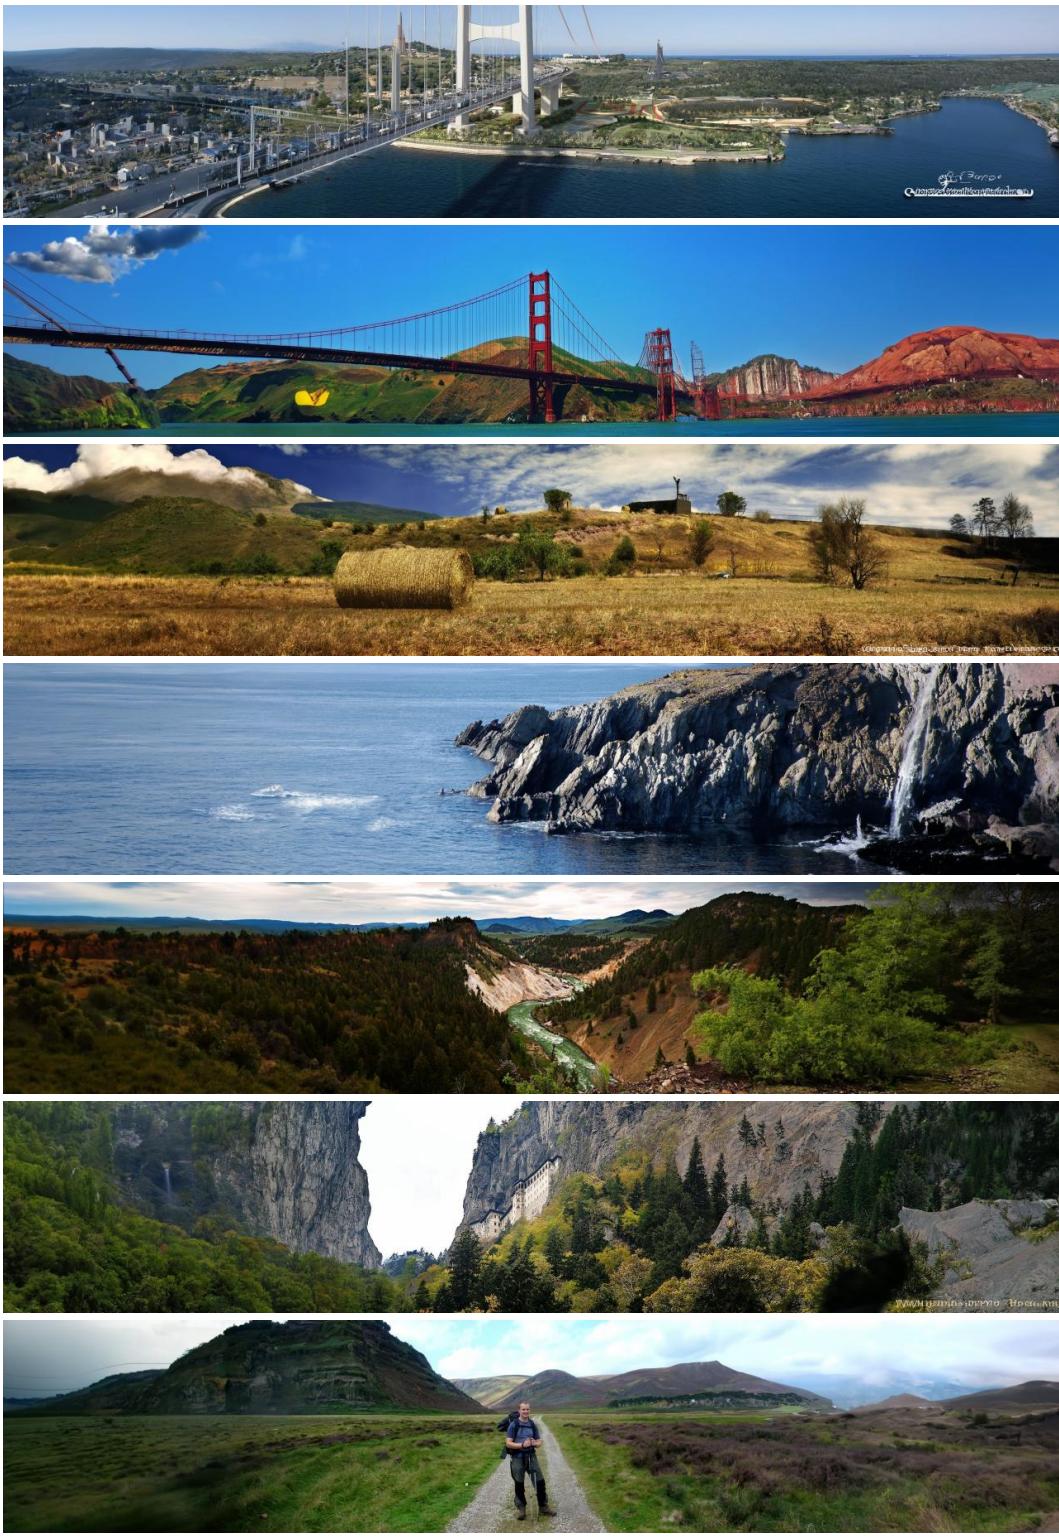


Figure C.10: Palette panorama uncropping. Given the center 256×256 pixels, we extrapolate 512 pixels to the right, in steps of 128 (via 50% uncropping tasks), and 512 pixels to the left, yielding a 256×1280 panorama.